

Art_Extract_Test1

March 29, 2025

```
[98]: # === IMPORTS ===
import os
import math
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.utils import class_weight
from tensorflow.keras import layers, models, losses, optimizers, applications, u
    mixed_precision
```

```
[184]: np.random.seed(42)
tf.random.set_seed(42)

# Configuration
BATCH_SIZE = 32
IMG_HEIGHT = 224
IMG_WIDTH = 224
NUM_EPOCHS = 7 # Increase epochs when in production
LEARNING_RATE = 0.001
TASK = "style" # Options: "style", "artist", "genre"
```

```
[101]: # Paths
BASE_DIR = "C:/Users/Ace/Gsoc_HumanAI/wikiart_csv"
WIKIART_DIR = "C:/Users/Ace/Gsoc_HumanAI/wikiart" # artwork images here
MODELS_DIR = "C:/Users/Ace/Gsoc_HumanAI"
TRAIN_DATA_PATH = f"{BASE_DIR}/{TASK}_train.csv"
VAL_DATA_PATH = f"{BASE_DIR}/{TASK}_val.csv"
CLASSES_PATH = f"{BASE_DIR}/{TASK}_class.txt"
```

```
[104]: from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight

def load_data(data_path, subset_size=1.0, random_state=42):
    """Load data from CSV file and apply stratified sampling."""
    df = pd.read_csv(data_path)
    df.columns = ['image_path', 'label']
```

```

    df['image_path'] = df['image_path'].apply(lambda x: os.path.
↪join(WIKIART_DIR, x))

    df = df[df['image_path'].apply(os.path.exists)]

    if subset_size < 1.0:
        df = df.groupby('label', group_keys=False).apply(
            lambda x: x.sample(frac=subset_size, random_state=random_state)
        )

    # Print sample paths for verification
    sample_paths = df['image_path'].sample(min(5, len(df))).tolist()
    for path in sample_paths:
        print(f"Checking if path exists: {path}")
        print(f"Exists: {os.path.exists(path)}")

    return df

def load_classes(classes_path):
    """Load class names from text file."""
    with open(classes_path, 'r') as f:
        classes = [line.strip() for line in f.readlines()]
    return classes

[106]: def preprocess_data(train_df, val_df, classes):
    """Preprocess data for training."""
    is_numeric_labels = isinstance(train_df['label'].iloc[0], (int, np.integer))

    if is_numeric_labels:
        train_df['label_encoded'] = train_df['label']
        val_df['label_encoded'] = val_df['label']

        label_map = {i: class_name for i, class_name in enumerate(classes)}

        train_df['label_name'] = train_df['label'].map(label_map)
        val_df['label_name'] = val_df['label'].map(label_map)

        le = LabelEncoder()
        le.fit(classes)
    else:
        le = LabelEncoder()
        le.fit(classes)

    unknown_train_labels = set(train_df['label']) - set(classes)
    unknown_val_labels = set(val_df['label']) - set(classes)

```

```

    if unknown_train_labels:
        print(f"Warning: Found {len(unknown_train_labels)} unknown labels_
↳in training data")
        print(f"Sample unknown labels: {list(unknown_train_labels)[:5]}")

        train_df = train_df[train_df['label'].isin(classes)]

    if unknown_val_labels:
        print(f"Warning: Found {len(unknown_val_labels)} unknown labels in_
↳validation data")
        print(f"Sample unknown labels: {list(unknown_val_labels)[:5]}")

        val_df = val_df[val_df['label'].isin(classes)]

    train_df['label_encoded'] = le.transform(train_df['label'])
    val_df['label_encoded'] = le.transform(val_df['label'])

    class_weights = compute_class_weight(
        'balanced',
        classes=np.unique(train_df['label_encoded']),
        y=train_df['label_encoded']
    )
    class_weights_dict = {i: weight for i, weight in enumerate(class_weights)}

    return train_df, val_df, le, class_weights_dict

```

```

[108]: def create_data_generators(df, batch_size, task):
    """Create a tf.data.Dataset compatible with EfficientNetB3."""
    def load_and_preprocess_image(path):
        image = tf.io.read_file(path)
        image = tf.image.decode_jpeg(image, channels=3)
        image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])
        image = tf.keras.applications.efficientnet.preprocess_input(image)
        return image

    def augment_image(image, label):
        if task == "style":
            image = tf.cast(image, tf.uint8)
            image = tf.image.random_flip_left_right(image)

            image = tf.image.random_brightness(image, max_delta=0.1)
            image = tf.image.random_contrast(image, lower=0.9, upper=1.1)
            image = tf.image.random_saturation(image, lower=0.9, upper=1.1)

            crop_factor = tf.random.uniform([], 0.9, 1.0)
            crop_size = tf.cast(
                tf.cast([IMG_HEIGHT, IMG_WIDTH], tf.float32) * crop_factor,

```

```

        tf.int32
    )
    crop_size = tf.minimum(crop_size, [IMG_HEIGHT, IMG_WIDTH])
    image = tf.image.random_crop(image, [crop_size[0], crop_size[1], 3])
    image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])

    elif task == "artist":
        image = tf.image.random_flip_left_right(image)
        image = tf.image.random_brightness(image, max_delta=0.2)
        image = tf.image.random_contrast(image, lower=0.8, upper=1.2)
        image = tf.image.random_saturation(image, lower=0.8, upper=1.2)

    elif task == "genre":
        image = tf.image.random_flip_left_right(image)

        image = tf.image.random_brightness(image, max_delta=0.1)
        image = tf.image.random_contrast(image, lower=0.9, upper=1.1)
        image = tf.image.random_saturation(image, lower=0.9, upper=1.1)

    if tf.random.uniform([], 0, 1) > 0.5:
        crop_factor = tf.random.uniform([], 0.95, 1.0)
        crop_size = tf.cast(
            tf.cast([IMG_HEIGHT, IMG_WIDTH], tf.float32) * crop_factor,
            tf.int32
        )
        image = tf.image.random_crop(image, [crop_size[0],
↵crop_size[1], 3])
        image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH])

    return image, label

paths = df['image_path'].values
labels = df['label_encoded'].values

dataset = tf.data.Dataset.from_tensor_slices((paths, labels))

dataset = dataset.map(
    lambda path, label: (load_and_preprocess_image(path), label),
    num_parallel_calls=tf.data.AUTOTUNE
)

dataset = dataset.map(augment_image, num_parallel_calls=tf.data.AUTOTUNE)

dataset = dataset.shuffle(buffer_size=10000)

dataset = dataset.batch(batch_size).prefetch(tf.data.AUTOTUNE)

```

```
return dataset
```

```
[132]: from tensorflow.keras.applications import EfficientNetB0
from tensorflow.keras.layers import Input, Dense, Dropout, BatchNormalization,
↳ GlobalAveragePooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping,
↳ ModelCheckpoint

def focal_loss(gamma=2., alpha=0.25):
    def focal_loss_fixed(y_true, y_pred):
        epsilon = tf.keras.backend.epsilon()
        y_pred = tf.clip_by_value(y_pred, epsilon, 1. - epsilon)
        y_true = tf.one_hot(tf.cast(y_true, tf.int32), tf.shape(y_pred)[-1])
        alpha_t = y_true * alpha + (1 - y_true) * (1 - alpha)
        loss = -alpha_t * (y_true * tf.math.pow(1. - y_pred, gamma) * tf.math.
↳ log(y_pred))
        return tf.reduce_sum(loss, axis=-1)
    return focal_loss_fixed

class CosineAnnealingWithRestarts(tf.keras.optimizers.schedules.
↳ LearningRateSchedule):
    def __init__(self, initial_lr, T_max, eta_min=0):
        self.initial_lr = initial_lr
        self.T_max = T_max
        self.eta_min = eta_min
        self.t = 0

    def __call__(self, step):
        cos_inner = tf.math.pi * (self.t % self.T_max) / self.T_max
        lr = self.eta_min + (self.initial_lr - self.eta_min) * (1 + tf.math.
↳ cos(cos_inner)) / 2
        self.t += 1
        return lr

def build_conv_recurrent_model(num_classes, task, img_height=224,
↳ img_width=224):
    """Build improved models for art classification with fixes for artist task.
    """
    if task == "style" or task == "genre":
        base_model = tf.keras.applications.EfficientNetB2(
            weights='imagenet',
            include_top=False,
            input_shape=(img_height, img_width, 3)
        )
```

```

else:
    base_model = tf.keras.applications.EfficientNetB2(
        weights='imagenet',
        include_top=False,
        input_shape=(img_height, img_width, 3)
    )

if task == "style":
    for layer in base_model.layers[:-30]:
        layer.trainable = False
elif task == "genre":
    for layer in base_model.layers[:-60]:
        layer.trainable = False
else:
    for layer in base_model.layers[:-40]:
        layer.trainable = False

inputs = Input(shape=(img_height, img_width, 3))
x = tf.keras.applications.efficientnet.preprocess_input(inputs)
x = base_model(x)

x = GlobalAveragePooling2D()(x)

if task == "artist":
    x = Dense(1536, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.4)(x)

    shortcut =Dense(768)(x)

    x = Dense(768, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.3)(x)

    se = Dense(128, activation='relu')(x)
    se = Dense(768, activation='sigmoid')(se)
    x = x * se

    x = x + shortcut

elif task == "style":
    x = Dense(1024, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.3)(x)

    x = Dense(512, activation='relu')(x)

```

```

        x = BatchNormalization()(x)
        x = Dropout(0.3)(x)

        x = Dense(256, activation='relu')(x)
        x = BatchNormalization()(x)
        x = Dropout(0.2)(x)

    elif task == "genre":
        x = Dense(1024, activation='relu')(x)
        x = BatchNormalization()(x)
        x = Dropout(0.3)(x)

        x = Dense(512, activation='relu')(x)
        x = BatchNormalization()(x)
        x = Dropout(0.25)(x)

        x = Dense(256, activation='relu')(x)
        x = BatchNormalization()(x)
        x = Dropout(0.2)(x)

    outputs = Dense(num_classes, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    if task == "style":
        optimizer = Adam(learning_rate=CosineAnnealingWithRestarts(2e-3, 1000),
↪weight_decay=1e-5)
    elif task == "genre":
        initial_learning_rate = 5e-4
        lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
            initial_learning_rate,
            decay_steps=2000,
            decay_rate=0.95,
            staircase=True
        )
        optimizer = Adam(learning_rate=lr_schedule, weight_decay=1e-5)
    else:
        optimizer = Adam(learning_rate=CosineAnnealingWithRestarts(5e-4, 2000),
↪weight_decay=1e-5)

    model.compile(
        optimizer=optimizer,
        loss=focal_loss(gamma=2.0, alpha=0.25),
        metrics=['accuracy']
    )

    return model

```

```

[122]: import seaborn as sns
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report, f1_score
from tensorflow.keras.utils import plot_model

def evaluate_model(model, val_dataset, le, TASK):
    """Evaluate the model and handle both integer and one-hot encoded labels."""
    # Create results folder
    os.makedirs('results', exist_ok=True)

    plot_model(
        model,
        to_file=f'results/model_architecture_{TASK}.png',
        show_shapes=True,
        show_layer_names=True,
        expand_nested=True
    )
    print(f"Model architecture saved as 'results/model_architecture_{TASK}.png'")

    predictions = model.predict(val_dataset, steps=len(val_dataset), verbose=1)
    predicted_classes = np.argmax(predictions, axis=1)

    true_classes = []
    for _, labels in val_dataset:
        labels = labels.numpy()
        if labels.ndim == 2:
            true_classes.extend(np.argmax(labels, axis=1))
        else:
            true_classes.extend(labels)

    true_classes = np.array(true_classes)

    accuracy = accuracy_score(true_classes, predicted_classes)
    f1 = f1_score(true_classes, predicted_classes, average='weighted')

    print(f"\nAccuracy: {accuracy:.4f}")
    print(f"F1 Score: {f1:.4f}\n")

    conf_matrix = confusion_matrix(true_classes, predicted_classes)
    class_report = classification_report(true_classes, predicted_classes, \
        target_names=le.classes_)

    with open(f'results/classification_report_{TASK}.txt', 'w') as f:
        f.write(class_report)

```



```

plt.figure(figsize=(14, 12))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_,
            cbar_kws={'shrink': 0.8}, linewidths=0.5, linecolor='gray',
            annot_kws={"size": 7},
            vmin=0, vmax=conf_matrix.max())

plt.title(f'Confusion Matrix for {TASK}')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

plt.xticks(ticks=np.arange(len(le.classes_) + 0.5, labels=le.classes_,
↪rotation=45, ha='right', fontsize=8)
plt.yticks(ticks=np.arange(len(le.classes_) + 0.5, labels=le.classes_,
↪rotation=0, va='center', fontsize=8)

plt.subplots_adjust(left=0.3, bottom=0.2)

plt.savefig(f'results/confusion_matrix_{TASK}.png')
plt.show()

f1_scores = f1_score(true_classes, predicted_classes, average=None)
plt.figure(figsize=(16, 6))
sns.barplot(x=le.classes_, y=f1_scores, palette='viridis')

plt.title(f'Per-Class F1 Scores for {TASK}')
plt.ylabel('F1 Score')

plt.xticks(ticks=np.arange(len(le.classes_)), labels=le.classes_,
↪rotation=45, ha='right')

plt.tight_layout()
plt.savefig(f'results/f1_scores_{TASK}.png')
plt.show()

with open(f'results/class_accuracy_{TASK}.txt', 'w') as f:
    f.write("Class-wise Evaluation:\n")
    for i, class_name in enumerate(le.classes_):
        class_acc = conf_matrix[i, i] / conf_matrix[i].sum() if
↪conf_matrix[i].sum() > 0 else 0
        f.write(f"{class_name} - Accuracy: {class_acc:.4f}\n")

prediction_confidence = np.max(predictions, axis=1)
low_confidence_indices = np.where(prediction_confidence < 0.5)[0]
misclassified_indices = np.where(predicted_classes != true_classes)[0]

```

```

outlier_indices = np.union1d(low_confidence_indices, misclassified_indices)

print(f"\nFound {len(outlier_indices)} potential outliers")

return outlier_indices, predictions, true_classes, predicted_classes

```

```

[124]: def visualize_training_history(history, TASK):
        """Visualize training history"""
        plt.figure(figsize=(12, 4))

        plt.subplot(1, 2, 1)
        plt.plot(history['accuracy'], label='Train Accuracy')
        plt.plot(history['val_accuracy'], label='Validation Accuracy')
        plt.title('Model Accuracy')
        plt.ylabel('Accuracy')
        plt.xlabel('Epoch')
        plt.legend(loc='upper left')

        plt.subplot(1, 2, 2)
        plt.plot(history['loss'], label='Train Loss')
        plt.plot(history['val_loss'], label='Validation Loss')
        plt.title('Model Loss')
        plt.ylabel('Loss')
        plt.xlabel('Epoch')
        plt.legend(loc='upper left')

        plt.tight_layout()
        plt.savefig(f'training_history_{TASK}.png')
        plt.show()

def visualize_outliers(outlier_indices, val_dataset, predictions, true_classes,
    predicted_classes, le, task, num_examples=5):
    """Visualize outlier examples"""

    results_dir = f'results/{task}'
    os.makedirs(results_dir, exist_ok=True)

    images = []
    for img_batch, _ in val_dataset:
        images.extend(img_batch.numpy())

    if len(outlier_indices) > num_examples:
        sample_indices = np.random.choice(outlier_indices, num_examples,
            replace=False)
    else:

```

```

sample_indices = outlier_indices

plt.figure(figsize=(12, 12))
for i, idx in enumerate(sample_indices):
    if idx < len(images):
        image = images[idx]

        if image.max() <= 1.0:
            image = (image * 255).astype('uint8')
        else:
            image = np.clip(image, 0, 255).astype('uint8')

        true_label = le.classes_[true_classes[idx]]
        pred_label = le.classes_[predicted_classes[idx]]
        confidence = predictions[idx][predicted_classes[idx]]

        plt.subplot(3, 2, i + 1)
        plt.imshow(image)
        plt.title(f"True: {true_label}\nPred: {pred_label}\nConf: {confidence:.2f}",
            ↪fontsize=10)
        plt.axis('off')

plt.tight_layout()
outlier_path = os.path.join(results_dir, f'outliers_{task}.png')
plt.savefig(outlier_path)
plt.show()
print(f"Outliers saved to: {outlier_path}")

```

```

[126]: # Data exploration function
def explore_dataset():
    """Explore the dataset structure"""
    for task in ["artist", "genre", "style"]:
        train_path = f"{BASE_DIR}/{task}_train.csv"
        val_path = f"{BASE_DIR}/{task}_val.csv"
        class_path = f"{BASE_DIR}/{task}_class.txt"

        print(f"\n{'='*40}")
        print(f"Exploring {task.upper()} dataset")
        print(f"{'='*40}")

        print(f"Train file exists: {os.path.exists(train_path)}")
        print(f"Val file exists: {os.path.exists(val_path)}")
        print(f"Class file exists: {os.path.exists(class_path)}")

        try:
            train_df = pd.read_csv(train_path)
            val_df = pd.read_csv(val_path)

```

```

print(f"\nTrain data shape: {train_df.shape}")
print(f"Validation data shape: {val_df.shape}")

print(f"\nTrain columns: {train_df.columns.tolist()}")

print("\nSample train data (first 3 rows):")
print(train_df.head(3))

if len(train_df.columns) >= 2:
    label_col = train_df.iloc[:, 1]
    unique_labels = label_col.unique()
    print(f"\nNumber of unique labels in training data: {len(unique_labels)}")
    print(f"Sample labels: {unique_labels[:5].tolist()}")

if len(train_df.columns) >= 1:
    img_col = train_df.iloc[:, 0]
    sample_paths = img_col.sample(min(3, len(img_col))).tolist()
    print("\nSample image paths:")
    for path in sample_paths:
        print(f" {path}")
        full_path = os.path.join(WIKIART_DIR, path)
        print(f" Exists in wikiart folder: {os.path.exists(full_path)}")

except Exception as e:
    print(f"Error exploring {task} dataset: {str(e)}")

print("Exploring dataset structure...")
explore_dataset()

print("\nStarting model training...")

```

Exploring dataset structure...

```

=====
Exploring ARTIST dataset
=====

```

```

Train file exists: True
Val file exists: True
Class file exists: True

```

```

Train data shape: (13345, 2)
Validation data shape: (5705, 2)

```

```

Train columns: ['Realism/vincent-van-gogh_pine-trees-in-the-fen-1884.jpg', '22']

```

Sample train data (first 3 rows):

	Realism/vincent-van-gogh_pine-trees-in-the-fen-1884.jpg	22
0	Baroque/rembrandt_the-angel-appearing-to-the-s...	20
1	Post_Impressionism/paul-cezanne_portrait-of-th...	16
2	Impressionism/pierre-auguste-renoir_young-girl...	17

Number of unique labels in training data: 23

Sample labels: [20, 16, 17, 9, 1]

Sample image paths:

Romanticism/gustave-dore_paradise-lost-4.jpg
Exists in wikiart folder: True
Post_Impressionism/vincent-van-gogh_the-raising-of-lazarus-1890.jpg
Exists in wikiart folder: True
Impressionism/eugene-boudin_the-port-of-deauville-1.jpg
Exists in wikiart folder: True

=====

Exploring GENRE dataset

=====

Train file exists: True

Val file exists: True

Class file exists: True

Train data shape: (45502, 2)

Validation data shape: (19491, 2)

Train columns: ['Northern_Renaissance/hieronymus-bosch_st-jacques-and-the-magician-hermogenes.jpg', '7']

Sample train data (first 3 rows):

	Northern_Renaissance/hieronymus-bosch_st-jacques-and-the-magician-hermogenes.jpg	\
0	Post_Impressionism/vincent-van-gogh_ears-of-wh...	
1	Symbolism/theodor-severin-kittelsen_kvitebj-rn...	
2	Expressionism/martiros-saryan_mother-of-the-ar...	

	7
0	4
1	3
2	6

Number of unique labels in training data: 10

Sample labels: [4, 3, 6, 8, 0]

Sample image paths:

Post_Impressionism/wassily-kandinsky_gabriele-munter-1905.jpg

```
Exists in wikiart folder: True
Impressionism/childe-hassam_clarissa.jpg
Exists in wikiart folder: True
Post_Impressionism/pyotr-konchalovsky_pine-tree-1921.jpg
Exists in wikiart folder: True
```

```
=====
Exploring STYLE dataset
=====
```

```
Train file exists: True
Val file exists: True
Class file exists: True
```

```
Train data shape: (57024, 2)
Validation data shape: (24420, 2)
```

```
Train columns: ['Impressionism/edgar-degas_landscape-on-the-orne.jpg', '12']
```

```
Sample train data (first 3 rows):
```

```
Impressionism/edgar-degas_landscape-on-the-orne.jpg 12
0 Realism/camille-corot_mantes-cathedral.jpg 21
1 Abstract_Expressionism/gene-davis_untitled-197... 0
2 Symbolism/kuzma-petrov-vodkin_in-the-1920.jpg 24
```

```
Number of unique labels in training data: 27
Sample labels: [21, 0, 24, 12, 7]
```

```
Sample image paths:
```

```
Art_Nouveau_Modern/felix-vallotton_portrait-of-belgian-symbolist-poet-max-
elskamp-1898.jpg
Exists in wikiart folder: True
Impressionism/edmund-charles-tarbell_the-bath-1893.jpg
Exists in wikiart folder: True
Impressionism/pierre-auguste-renoir_rocks-with-shrimp-fishermen-1892.jpg
Exists in wikiart folder: True
```

```
Starting model training...
```

```
[186]: def train_model(task, subset_size=1.0):
        """Train the improved model with enhanced training strategy."""
        train_df = load_data(f"{BASE_DIR}/{task}_train.csv",
        ↪subset_size=subset_size)
        val_df = load_data(f"{BASE_DIR}/{task}_val.csv", subset_size=subset_size)
        classes = load_classes(f"{BASE_DIR}/{task}_class.txt")

        train_df, val_df, label_encoder, class_weights_dict =
        ↪preprocess_data(train_df, val_df, classes)
```

```

train_dataset = create_data_generators(train_df, BATCH_SIZE, task)
val_dataset = create_data_generators(val_df, BATCH_SIZE, task)

model = build_conv_recurrent_model(len(classes), task)
print(model.summary())

# Callbacks for better training
checkpoint = ModelCheckpoint(
    f'{MODELS_DIR}/models/best_model_{task}.keras',
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

early_stopping = EarlyStopping(
    monitor='val_accuracy',
    patience=12,
    restore_best_weights=True,
    verbose=1
)

reduce_lr = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=4,
    min_lr=1e-7,
    verbose=1
)

model.compile(
    optimizer=Adam(learning_rate=LEARNING_RATE, weight_decay=1e-6),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

print("Phase 1: Initial training with mostly frozen base model...")
history1 = model.fit(
    train_dataset,
    epochs=NUM_EPOCHS,
    validation_data=val_dataset,
    callbacks=[checkpoint, early_stopping, reduce_lr],
    class_weight=class_weights_dict
)

print("Phase 2: Fine-tuning with more layers unfrozen...")

```

```

base_model = model.layers[1]

for layer in base_model.layers:
    layer.trainable = True

model.compile(
    optimizer=Adam(learning_rate=LEARNING_RATE/10, weight_decay=1e-6),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

history2 = model.fit(
    train_dataset,
    epochs=10,
    initial_epoch=history1.epoch[-1] + 1,
    validation_data=val_dataset,
    callbacks=[checkpoint, early_stopping, reduce_lr],
    class_weight=class_weights_dict
)

# Evaluate the model
print("Evaluating model...")
outlier_indices, predictions, true_classes, predicted_classes = evaluate_model(
    model, val_dataset, label_encoder, task
)

print("Visualizing results...")
combined_history = {}
for k in history1.history.keys():
    combined_history[k] = history1.history[k] + history2.history[k]

visualize_training_history(combined_history, task)
visualize_outliers(
    outlier_indices, val_dataset, predictions,
    true_classes, predicted_classes, label_encoder, task
)

return model, combined_history

```

```

[ ]: def train_all_models_improved():
    """Train improved models for all tasks with stratified sampling."""
    results = {}
    subset_sizes = {
        "artist": 1.0, # Use 100% of the dataset for Artist
        "genre": 0.5, # Use 50% of the dataset for Genre
    }

```



```

        "style": 0.3    # Use 30% of the dataset for Style
    }

    for task in ["style", "genre", "artist"]:
        print(f"\n\n{'='*60}")
        print(f"Training improved model for {task.upper()}")
        print(f"{'='*60}\n")

        model, history = train_model(task, subset_size=subset_sizes[task])
        results[task] = (model, history)

    return results

if __name__ == "__main__":
    # Training all models
    all_results = train_all_models_improved()

```

```

[164]: from tensorflow.keras.models import load_model
       from tensorflow.keras.preprocessing import image

```

```

[166]: def load_trained_model(task):
        """Load a trained model for a specific task"""
        model_path = f'{MODELS_DIR}/models/best_model_{task}.keras'
        if not os.path.exists(model_path):
            print(f"Error: Model for {task} not found at {model_path}")
            return None

        try:
            model = load_model(model_path)
            print(f"Successfully loaded {task} model")
            return model
        except Exception as e:
            print(f"Error loading {task} model: {str(e)}")
            return None

    def load_class_labels(task):
        """Load class labels for a specific task"""
        classes_path = f'{BASE_DIR}/{task}_class.txt'
        try:
            with open(classes_path, 'r') as f:
                classes = [line.strip() for line in f.readlines()]
            print(f"Loaded {len(classes)} {task} classes")
            return classes
        except Exception as e:
            print(f"Error loading {task} classes: {str(e)}")
            return None

```

```
[168]: def preprocess_image(img_path):
        """Preprocess an image for model prediction"""
        try:
            if not os.path.exists(img_path):
                print(f"Error: Image not found at {img_path}")
                return None

            # Load and preprocess the image
            img = image.load_img(img_path, target_size=(224, 224))
            img_array = image.img_to_array(img)
            img_array = np.expand_dims(img_array, axis=0)
            img_array = img_array / 255.0

            return img_array, img
        except Exception as e:
            print(f"Error preprocessing image: {str(e)}")
            return None, None

[170]: def predict_artwork(img_path, tasks=None):
        """
        Predict artist, style, and genre for a given artwork image

        Parameters:
        img_path (str): Path to the artwork image
        tasks (list): List of tasks to perform, default ["artist", "genre", "style"]

        Returns:
        dict: Dictionary with predictions for each task
        """
        if tasks is None:
            tasks = ["artist", "genre", "style"]

        img_array, original_img = preprocess_image(img_path)
        if img_array is None:
            return None

        results = {}

        models = {task: load_trained_model(task) for task in tasks}
        class_labels = {task: load_class_labels(task) for task in tasks}

        for task in tasks:
            print(f"\nPredicting {task}...")

            model = models.get(task)
            classes = class_labels.get(task)
```

```

        if model is None or classes is None:
            results[task] = {"error": f"Could not load model or classes for {task}"}
            continue

        try:
            img_tensor = tf.convert_to_tensor(img_array)
            img_tensor = tf.ensure_shape(img_tensor, (1, 224, 224, 3)) #
            Example shape (adjust to your model)

            @tf.function(reduce_retracing=True)
            def predict_step(input_tensor):
                return model(input_tensor)

            predictions = predict_step(img_tensor)

            top_indices = tf.argsort(predictions[0], direction="DESCENDING")[:3]
            top_predictions = [(classes[i.numpy()], float(predictions[0][i].
            numpy())) for i in top_indices]

            results[task] = {
                "top_predictions": top_predictions,
                "prediction": classes[int(tf.argmax(predictions[0]))],
                "confidence": float(tf.reduce_max(predictions[0]))
            }

            print(f"Top {task} predictions:")
            for class_name, prob in top_predictions:
                print(f"  {class_name}: {prob:.4f}")

        except Exception as e:
            print(f"Error making prediction for {task}: {str(e)}")
            results[task] = {"error": str(e)}

    visualize_prediction_results(img_path, original_img, results)

    return results

```

```

[172]: def visualize_prediction_results(img_path, original_img, results):
        """Visualize the prediction results"""
        plt.figure(figsize=(12, 8))

        plt.subplot(1, 2, 1)
        plt.imshow(original_img)
        plt.title(f"Artwork: {os.path.basename(img_path)}")
        plt.axis('off')

```

```

plt.subplot(1, 2, 2)
plt.axis('off')

result_text = "Prediction Results:\n\n"

for task in results:
    result_text += f"{task.capitalize()}\n"

    if "error" in results[task]:
        result_text += f"  Error: {results[task]['error']}\n"
    else:
        for i, (class_name, prob) in enumerate(results[task]["top_predictions"]):
            result_text += f"    {i+1}. {class_name}: {prob:.2%}\n"

        result_text += "\n"

plt.text(0.1, 0.5, result_text, fontsize=12, verticalalignment='center')

plt.tight_layout()
plt.savefig('artwork_prediction.png')
plt.show()

```

```

[174]: def batch_predict_artworks(folder_path, tasks=None):
    """
    Predict artist, style, and genre for all artwork images in a folder

    Parameters:
    folder_path (str): Path to the folder containing artwork images
    tasks (list): List of tasks to perform, default ["artist", "genre", "style"]
    """
    if tasks is None:
        tasks = ["artist", "genre", "style"]

    if not os.path.exists(folder_path):
        print(f"Error: Folder not found at {folder_path}")
        return

    image_extensions = ['.jpg', '.jpeg', '.png']
    image_files = [f for f in os.listdir(folder_path)
                    if os.path.isfile(os.path.join(folder_path, f)) and
                    any(f.lower().endswith(ext) for ext in image_extensions)]

    if not image_files:
        print(f"No image files found in {folder_path}")
        return

```

```

print(f"Found {len(image_files)} image files. Starting batch prediction...")

models = {}
class_labels = {}

for task in tasks:
    models[task] = load_trained_model(task)
    class_labels[task] = load_class_labels(task)

    if models[task] is None or class_labels[task] is None:
        print(f"Warning: Could not load model or classes for {task}")

all_results = {}
for img_file in image_files:
    img_path = os.path.join(folder_path, img_file)
    print(f"\nProcessing {img_file}...")

    img_array, _ = preprocess_image(img_path)
    if img_array is None:
        all_results[img_file] = {"error": "Failed to preprocess image"}
        continue

    img_results = {}
    for task in tasks:
        if models[task] is None or class_labels[task] is None:
            img_results[task] = {"error": f"Model or classes not available_
↳for {task}"}
            continue

        try:
            predictions = models[task].predict(img_array)

            top_indices = np.argsort(predictions[0])[-3:][::-1]
            top_predictions = [(class_labels[task][i],
↳float(predictions[0][i])) for i in top_indices]

            img_results[task] = {
                "top_predictions": top_predictions,
                "prediction": class_labels[task][np.argmax(predictions)],
                "confidence": float(np.max(predictions))
            }

        except Exception as e:
            print(f"Error making {task} prediction for {img_file}:_
↳{str(e)}")
            img_results[task] = {"error": str(e)}

```

```

        all_results[img_file] = img_results

    export_results_to_csv(all_results, folder_path)

    print(f"\nCompleted batch prediction for {len(image_files)} images")
    return all_results

```

```

[176]: def export_results_to_csv(all_results, folder_path):
        """Export batch prediction results to CSV"""
        import pandas as pd

        rows = []
        for img_file, img_results in all_results.items():
            row = {'image': img_file}

            for task in img_results:
                if "error" in img_results[task]:
                    row[f'{task}_prediction'] = "ERROR"
                    row[f'{task}_confidence'] = 0.0
                else:
                    row[f'{task}_prediction'] = img_results[task]["prediction"]
                    row[f'{task}_confidence'] = img_results[task]["confidence"]

                    # Add top 3 predictions
                    for i, (class_name, prob) in enumerate(
                        img_results[task]["top_predictions"]):
                        row[f'{task}_top{i+1}'] = class_name
                        row[f'{task}_top{i+1}_confidence'] = prob

            rows.append(row)

        df = pd.DataFrame(rows)
        csv_path = os.path.join(folder_path, 'artwork_predictions.csv')
        df.to_csv(csv_path, index=False)
        print(f"Results exported to {csv_path}")

    def analyze_single_image(img_path):
        """
        Analyze a single artwork image for artist, style, and genre

        Parameters:
        img_path (str): Path to the artwork image
        """
        print(f"Analyzing artwork: {img_path}")
        results = predict_artwork(img_path)

        if results:

```

```

        print("\nSummary of predictions:")
        for task, task_results in results.items():
            if "error" in task_results:
                print(f" {task.capitalize()}: Error - {task_results['error']}")
            else:
                print(f" {task.capitalize()}: {task_results['prediction']}_{
↪({task_results['confidence']:.2%})")

    return results

```

```

[ ]: if __name__ == "__main__":
    # Single image prediction
    print("\n==== Single Image Prediction =====")
    analyze_single_image(test_image_path)

    # Batch prediction example - uncomment to use
    # test_folder_path = "./test_images" # Change this to your test folder path
    # print("\n==== Batch Prediction =====")
    # batch_results = batch_predict_artworks(test_folder_path)

```

```

[ ]:

```

```
[150]: # Example usage
if __name__ == "__main__":

    # Single image prediction
    print("\n==== Single Image Prediction =====")
    analyze_single_image(test_image_path)

    # Batch prediction example - uncomment to use
    # test_folder_path = "./test_images" # Change this to your test folder path
    # print("\n==== Batch Prediction =====")
    # batch_results = batch_predict_artworks(test_folder_path)
```

```
./test_artwork.jpg C:/Users/Ace/Gsoc_HumanAI/wikiart/Analytical_Cubism/pablo-
picasso_woman-with-a-mandolin-1909.jpg
```

```
==== Single Image Prediction =====
```

```
Analyzing artwork: C:/Users/Ace/Gsoc_HumanAI/wikiart/Analytical_Cubism/pablo-
picasso_woman-with-a-mandolin-1909.jpg
```

```
Successfully loaded artist model
```

```
Successfully loaded genre model
```

```
Successfully loaded style model
```

```
Loaded 23 artist classes
```

```
Loaded 10 genre classes
```

```
Loaded 27 style classes
```

```
Predicting artist...
```

```
Top artist predictions:
```

```
6 Eugene_Boudin: 0.3835
```

```
0 Albrecht_Durer: 0.2899
```

```
20 Rembrandt: 0.2002
```

```
Predicting genre...
```

```
WARNING:tensorflow:5 out of the last 5 calls to <function
```

```
predict_artwork.<locals>.predict_step at 0x0000029D2FD958A0> triggered
```

```
tf.function retracing. Tracing is expensive and the excessive number of tracings
```

```
could be due to (1) creating @tf.function repeatedly in a loop, (2) passing
```

```
tensors with different shapes, (3) passing Python objects instead of tensors.
```

```
For (1), please define your @tf.function outside of the loop. For (2),
```

```
@tf.function has reduce_retracing=True option that can avoid unnecessary
```

```
retracing. For (3), please refer to
```


https://www.tensorflow.org/guide/function#controlling_retracing and
https://www.tensorflow.org/api_docs/python/tf/function for more details.

Top genre predictions:

0 abstract_painting: 0.9786
3 illustration: 0.0099
8 sketch_and_study: 0.0065

Predicting style...

WARNING:tensorflow:6 out of the last 6 calls to <function predict_artwork.<locals>.predict_step at 0x0000029D3489C180> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to

https://www.tensorflow.org/guide/function#controlling_retracing and
https://www.tensorflow.org/api_docs/python/tf/function for more details.

Top style predictions:

14 Minimalism: 0.6165
5 Color_Field_Painting: 0.3304
0 Abstract_Expressionism: 0.0426

Artwork: pablo-picasso_woman-with-a-mandolin-1909.jpg



Prediction Results:

Artist:

1. 6 Eugene_Boudin: 38.35%
2. 0 Albrecht_Durer: 28.99%
3. 20 Rembrandt: 20.02%

Genre:

1. 0 abstract_painting: 97.86%
2. 3 illustration: 0.99%
3. 8 sketch_and_study: 0.65%

Style:

1. 14 Minimalism: 61.65%
2. 5 Color_Field_Painting: 33.04%
3. 0 Abstract_Expressionism: 4.26%

Summary of predictions:

Artist: 6 Eugene_Boudin (38.35%)

Genre: 0 abstract_painting (97.86%)

Style: 14 Minimalism (61.65%)

[]:

```

        "style": 0.3    # Use 30% of the dataset for Style
    }

    for task in ["style", "genre", "artist"]:
        print(f"\n\n{'='*60}")
        print(f"Training improved model for {task.upper()}")
        print(f"{'='*60}\n")

        model, history = train_model(task, subset_size=subset_sizes[task])
        results[task] = (model, history)

    return results

if __name__ == "__main__":
    # Training all models
    all_results = train_all_models_improved()

```

```

=====
Training improved model for STYLE
=====

```

```

C:\Users\Ace\AppData\Local\Temp\ipykernel_37128\2906373448.py:14:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
    df = df.groupby('label', group_keys=False).apply(

Checking if path exists:
C:/Users/Ace/Gsoc_HumanAI/wikiart\Baroque/rembrandt_young-woman-with-a-broom.jpg
Exists: True
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Romanticism/cornelis-
springer_coming-out-of-church.jpg
Exists: True
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Realism/salvador-
dali_reclining-girl-in-sheep.jpg
Exists: True
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Expressionism/laszlo-
moholy-nagy_self-portrait-1919.jpg
Exists: True
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Realism/ivan-
shishkin_forest-into-the-frost.jpg
Exists: True

```

C:\Users\Ace\AppData\Local\Temp\ipykernel_37128\2906373448.py:14:
 DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
 This behavior is deprecated, and in a future version of pandas the grouping
 columns will be excluded from the operation. Either pass `include_groups=False`
 to exclude the groupings or explicitly select the grouping columns after groupby
 to silence this warning.

```
df = df.groupby('label', group_keys=False).apply(
```

Checking if path exists:

C:/Users/Ace/Gsoc_HumanAI/wikiart\Post_Impressionism/rene-magritte_the-staging-
 post-1948(1).jpg

Exists: True

Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Rococo/maurice-
 quentin-de-la-tour_face-of-the-man-after-alexis-grimou.jpg

Exists: True

Checking if path exists:

C:/Users/Ace/Gsoc_HumanAI/wikiart\Color_Field_Painting/ellsworth-
 kelly_rectangle-from-the-series-line-form-color-1951.jpg

Exists: True

Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Realism/john-singer-
 sargent_mrs-frederick-mead-mary-eliza-scribner.jpg

Exists: True

Checking if path exists:

C:/Users/Ace/Gsoc_HumanAI/wikiart\Post_Impressionism/jan-sluyters_the-white-
 tree.jpg

Exists: True

Model: "functional_5"

Layer (type)	Output Shape	Param #
input_layer_11 (InputLayer)	(None , 224, 224, 3)	0
efficientnetb2 (Functional)	(None , 7, 7, 1408)	7,768,569
global_average_pooling2d_5 (GlobalAveragePooling2D)	(None , 1408)	0
dense_20 (Dense)	(None , 1024)	1,442,816
batch_normalization_15 (BatchNormalization)	(None , 1024)	4,096
dropout_15 (Dropout)	(None , 1024)	0
dense_21 (Dense)	(None , 512)	524,800

batch_normalization_16 (BatchNormalization)	(None, 512)	2,048
dropout_16 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 256)	131,328
batch_normalization_17 (BatchNormalization)	(None, 256)	1,024
dropout_17 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 27)	6,939

Total params: 9,881,620 (37.70 MB)

Trainable params: 5,084,135 (19.39 MB)

Non-trainable params: 4,797,485 (18.30 MB)

None

Phase 1: Initial training with mostly frozen base model...

Epoch 1/7

535/535 0s 1s/step -

accuracy: 0.1519 - loss: 3.2685

Epoch 1: val_accuracy improved from -inf to 0.23492, saving model to

C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras

535/535 984s 2s/step -

accuracy: 0.1520 - loss: 3.2675 - val_accuracy: 0.2349 - val_loss: 2.6823 -

learning_rate: 0.0010

Epoch 2/7

535/535 0s 1s/step -

accuracy: 0.2718 - loss: 2.1875

Epoch 2: val_accuracy improved from 0.23492 to 0.29252, saving model to

C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras

535/535 898s 2s/step -

accuracy: 0.2718 - loss: 2.1872 - val_accuracy: 0.2925 - val_loss: 2.2492 -

learning_rate: 0.0010

Epoch 3/7

535/535 0s 1s/step -

accuracy: 0.3474 - loss: 1.8038

Epoch 3: val_accuracy improved from 0.29252 to 0.30713, saving model to

C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras

535/535 888s 2s/step -

accuracy: 0.3474 - loss: 1.8036 - val_accuracy: 0.3071 - val_loss: 2.3143 -

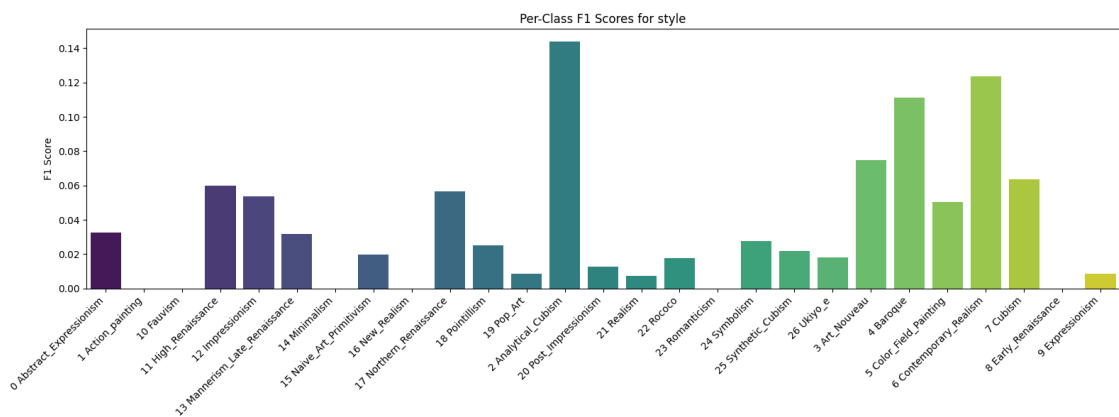

```

learning_rate: 0.0010
Epoch 4/7
535/535          0s 1s/step -
accuracy: 0.3711 - loss: 1.6134
Epoch 4: val_accuracy improved from 0.30713 to 0.35449, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          899s 2s/step -
accuracy: 0.3710 - loss: 1.6133 - val_accuracy: 0.3545 - val_loss: 2.1200 -
learning_rate: 0.0010
Epoch 5/7
535/535          0s 1s/step -
accuracy: 0.3985 - loss: 1.4677
Epoch 5: val_accuracy improved from 0.35449 to 0.35613, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          888s 2s/step -
accuracy: 0.3985 - loss: 1.4676 - val_accuracy: 0.3561 - val_loss: 2.2112 -
learning_rate: 0.0010
Epoch 6/7
535/535          0s 1s/step -
accuracy: 0.4395 - loss: 1.4128
Epoch 6: val_accuracy improved from 0.35613 to 0.37374, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          889s 2s/step -
accuracy: 0.4394 - loss: 1.4127 - val_accuracy: 0.3737 - val_loss: 2.0529 -
learning_rate: 0.0010
Epoch 7/7
535/535          0s 1s/step -
accuracy: 0.4591 - loss: 1.2794
Epoch 7: val_accuracy improved from 0.37374 to 0.37906, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          882s 2s/step -
accuracy: 0.4590 - loss: 1.2794 - val_accuracy: 0.3791 - val_loss: 2.0281 -
learning_rate: 0.0010
Restoring model weights from the end of the best epoch: 7.
Phase 2: Fine-tuning with more layers unfrozen...
Epoch 8/10
535/535          0s 5s/step -
accuracy: 0.4083 - loss: 1.3818
Epoch 8: val_accuracy improved from 0.37906 to 0.42219, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          3269s 6s/step -
accuracy: 0.4083 - loss: 1.3814 - val_accuracy: 0.4222 - val_loss: 1.7743 -
learning_rate: 1.0000e-04
Epoch 9/10
535/535          0s 5s/step -
accuracy: 0.4902 - loss: 1.0179
Epoch 9: val_accuracy improved from 0.42219 to 0.44458, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras


```

```
535/535          3135s 6s/step -
accuracy: 0.4902 - loss: 1.0178 - val_accuracy: 0.4446 - val_loss: 1.7096 -
learning_rate: 1.0000e-04
Epoch 10/10
535/535          0s 6s/step -
accuracy: 0.5310 - loss: 0.8715
Epoch 10: val_accuracy improved from 0.44458 to 0.46178, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_style.keras
535/535          3479s 6s/step -
accuracy: 0.5310 - loss: 0.8714 - val_accuracy: 0.4618 - val_loss: 1.6927 -
learning_rate: 1.0000e-04
Restoring model weights from the end of the best epoch: 10.
Evaluating model...
Model architecture saved as 'results/model_architecture_style.png'
229/229          240s 879ms/step

Accuracy: 0.0700
F1 Score: 0.0749
```



Found 7107 potential outliers
Visualizing results...



True: 13 Mannerism_Late_Renaissance
 Pred: 0 Abstract_Expressionism
 Conf: 0.42



True: 6 Contemporary_Realism
 Pred: 4 Baroque
 Conf: 0.47



True: 4 Baroque
 Pred: 7 Cubism
 Conf: 0.46



True: 2 Analytical_Cubism
 Pred: 3 Art_Nouveau
 Conf: 0.53



True: 3 Art_Nouveau
 Pred: 11 High_Renaissance
 Conf: 0.27



Outliers saved to: results/style/outliers_style.png

```
=====
Training improved model for GENRE
=====
```

```
C:\Users\Ace\AppData\Local\Temp\ipykernel_37128\2906373448.py:14:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
```

```
df = df.groupby('label', group_keys=False).apply(
```

```
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Rococo/thomas-
gainsborough_portrait-of-george-iii-1781.jpg
```

```
Exists: True
```

```
Checking if path exists:
```

```
C:/Users/Ace/Gsoc_HumanAI/wikiart\Post_Impressionism/vincent-van-gogh_the-
bridge.jpg
```

```
Exists: True
```

```
Checking if path exists:
```

```
C:/Users/Ace/Gsoc_HumanAI/wikiart\Northern_Renaissance/lucas-cranach-the-
elder_christ-carrying-the-cross-1538.jpg
```

```
Exists: True
```

```
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Romanticism/jan-
matejko_sigmund-and-barbara.jpg
```

```
Exists: True
```

```
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Symbolism/kuzma-
petrov-vodkin_portrait-of-ria-portrait-of-a-a-kholopova-1915.jpg
```

```
Exists: True
```

```
C:\Users\Ace\AppData\Local\Temp\ipykernel_37128\2906373448.py:14:
```

```
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
```

```
df = df.groupby('label', group_keys=False).apply(
```

```
Checking if path exists:
```

```
C:/Users/Ace/Gsoc_HumanAI/wikiart\High_Renaissance/vittore-carpaccio_portrait-
of-a-lady-1.jpg
```

```
Exists: True
```

```
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Realism/ivan-
shishkin_pines-sunny-day.jpg
```

```
Exists: True
```

```
Checking if path exists:
```

```
C:/Users/Ace/Gsoc_HumanAI/wikiart\Post_Impressionism/georges-braque_nude.jpg
```

```
Exists: True
```

```
Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart\Realism/adolf-
```

hitler_the-st-charles-church.jpg

Exists: True

Checking if path exists: C:/Users/Ace/Gsoc_HumanAI/wikiart/Realism/winslow-homer_the-bridal-path-white-mountains-1868.jpg

Exists: True

Model: "functional_6"

Layer (type)	Output Shape	Param #
input_layer_13 (InputLayer)	(None , 224, 224, 3)	0
efficientnetb2 (Functional)	(None , 7, 7, 1408)	7,768,569
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None , 1408)	0
dense_24 (Dense)	(None , 1024)	1,442,816
batch_normalization_18 (BatchNormalization)	(None , 1024)	4,096
dropout_18 (Dropout)	(None , 1024)	0
dense_25 (Dense)	(None , 512)	524,800
batch_normalization_19 (BatchNormalization)	(None , 512)	2,048
dropout_19 (Dropout)	(None , 512)	0
dense_26 (Dense)	(None , 256)	131,328
batch_normalization_20 (BatchNormalization)	(None , 256)	1,024
dropout_20 (Dropout)	(None , 256)	0
dense_27 (Dense)	(None , 10)	2,570

Total params: 9,877,251 (37.68 MB)

Trainable params: 6,453,502 (24.62 MB)

Non-trainable params: 3,423,749 (13.06 MB)

None

Phase 1: Initial training with mostly frozen base model...

Epoch 1/7

711/711 0s 1s/step -

accuracy: 0.6132 - loss: 1.1923

Epoch 1: val_accuracy improved from -inf to 0.66383, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_genre.keras

711/711 1449s 2s/step -

accuracy: 0.6132 - loss: 1.1922 - val_accuracy: 0.6638 - val_loss: 1.0750 -
learning_rate: 0.0010

Epoch 2/7

711/711 0s 1s/step -

accuracy: 0.6929 - loss: 0.8629

Epoch 2: val_accuracy did not improve from 0.66383

711/711 1305s 2s/step -

accuracy: 0.6929 - loss: 0.8629 - val_accuracy: 0.6301 - val_loss: 1.1054 -
learning_rate: 0.0010

Epoch 3/7

711/711 0s 1s/step -

accuracy: 0.7378 - loss: 0.7055

Epoch 3: val_accuracy improved from 0.66383 to 0.69020, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_genre.keras

711/711 1321s 2s/step -

accuracy: 0.7378 - loss: 0.7055 - val_accuracy: 0.6902 - val_loss: 0.9710 -
learning_rate: 0.0010

Epoch 4/7

711/711 0s 1s/step -

accuracy: 0.7599 - loss: 0.6266

Epoch 4: val_accuracy did not improve from 0.69020

711/711 1303s 2s/step -

accuracy: 0.7599 - loss: 0.6267 - val_accuracy: 0.6814 - val_loss: 1.0094 -
learning_rate: 0.0010

Epoch 5/7

711/711 0s 1s/step -

accuracy: 0.7829 - loss: 0.5395

Epoch 5: val_accuracy improved from 0.69020 to 0.71288, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_genre.keras

711/711 1307s 2s/step -

accuracy: 0.7829 - loss: 0.5396 - val_accuracy: 0.7129 - val_loss: 0.9178 -
learning_rate: 0.0010

Epoch 6/7

711/711 0s 1s/step -

accuracy: 0.8027 - loss: 0.4495

Epoch 6: val_accuracy did not improve from 0.71288

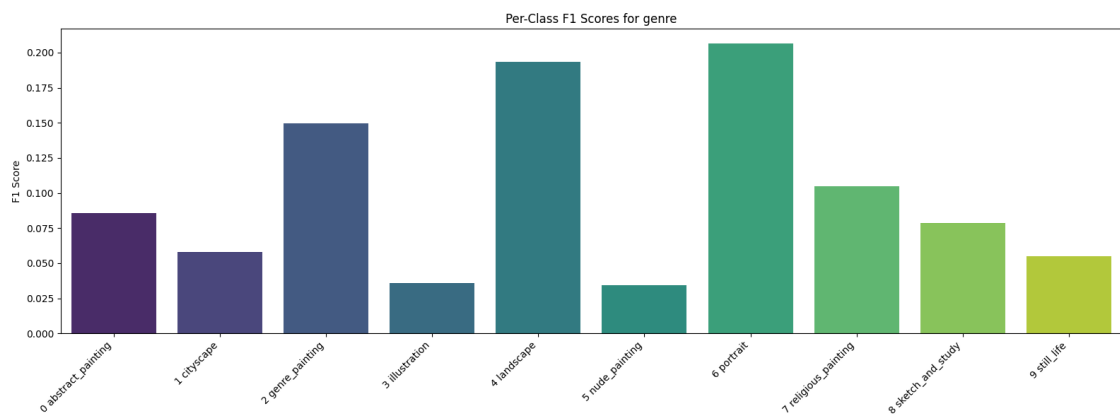

711/711 1299s 2s/step -

```


accuracy: 0.8027 - loss: 0.4496 - val_accuracy: 0.6908 - val_loss: 1.0042 -
learning_rate: 0.0010
Epoch 7/7
711/711          0s 1s/step -
accuracy: 0.8107 - loss: 0.4533
Epoch 7: val_accuracy did not improve from 0.71288
711/711          1314s 2s/step -
accuracy: 0.8107 - loss: 0.4533 - val_accuracy: 0.7056 - val_loss: 0.9698 -
learning_rate: 0.0010
Restoring model weights from the end of the best epoch: 5.
Phase 2: Fine-tuning with more layers unfrozen..
Epoch 8/10
711/711          0s 5s/step -
accuracy: 0.7141 - loss: 0.6954
Epoch 8: val_accuracy did not improve from 0.71288
711/711          4568s 6s/step -
accuracy: 0.7142 - loss: 0.6952 - val_accuracy: 0.7123 - val_loss: 0.8809 -
learning_rate: 1.0000e-04
Epoch 9/10
711/711          0s 5s/step -
accuracy: 0.7961 - loss: 0.4513
Epoch 9: val_accuracy improved from 0.71288 to 0.73484, saving model to
C:/Users/Ace/Gsoc_HumanAI/models/best_model_genre.keras
711/711          4495s 6s/step -
accuracy: 0.7962 - loss: 0.4512 - val_accuracy: 0.7348 - val_loss: 0.8303 -
learning_rate: 1.0000e-04
Epoch 10/10
711/711          0s 5s/step -
accuracy: 0.8349 - loss: 0.3517
Epoch 10: val_accuracy did not improve from 0.73484
711/711          4314s 6s/step -
accuracy: 0.8349 - loss: 0.3516 - val_accuracy: 0.7333 - val_loss: 0.8543 -
learning_rate: 1.0000e-04
Restoring model weights from the end of the best epoch: 9.
Evaluating model...
Model architecture saved as 'results/model_architecture_genre.png'
305/305          326s 908ms/step

Accuracy: 0.1356
F1 Score: 0.1401


```




Found 8521 potential outliers
Visualizing results...



True: 2 genre_painting
Pred: 1 cityscape
Conf: 0.99




True: 2 genre_painting
Pred: 9 still_life
Conf: 0.99



True: 0 abstract_painting
Pred: 4 landscape
Conf: 0.64



True: 2 genre_painting
Pred: 3 illustration
Conf: 0.93



True: 0 abstract_painting
Pred: 7 religious_painting
Conf: 0.70



Outliers saved to: results/genre/outliers_genre.png